

Traffic Anomaly Detection in SDN Using Deep Learning

Motivation:

- Rule-based firewalls are rigid and outdated
- SDN offers control, but opens new attack risks
- Deep Learning enables smart, real-time anomaly detection

Key Points:

- Simulate SDN using Mininet and Floodlight
- Collect flow-level traffic data using Floodlight Rest APIs
- Use DL models such as LSTM/Transformers to detect unusual traffic behaviour in real time

Simulation Topology & Floodlight Controller

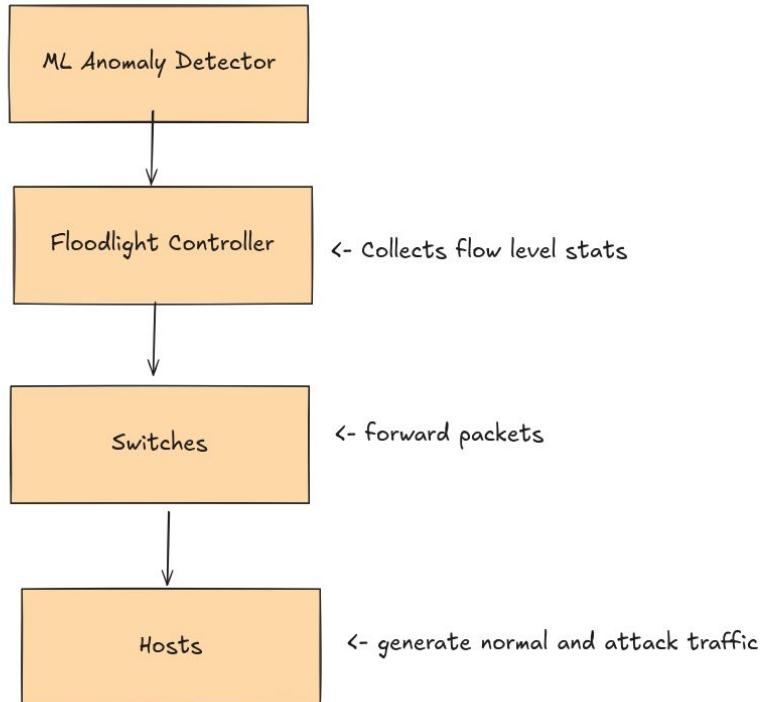
Network Setup:

- 6 Hosts (h1–h6)
- 3 Switches (s1–s3) in tree topology (depth=2, fanout=2)
- 1 Floodlight Controller running in Docker

Traffic:

- Normal: ping, iperf, curl
- Attacks: hping3, port scanning scripts

System Architecture



Data Collection

Network Topology:

- Mininet: 6 hosts, 3 OpenFlow switches
- Floodlight controller (OpenFlow 1.3)
- Tree topology

Traffic Generation:

Normal Traffic:

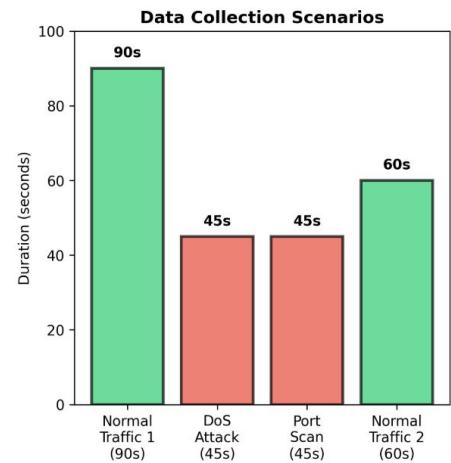
- ICMP (ping)
- TCP (iperf)
- HTTP (wget)

Attack Traffic:

- DoS attacks (flooding)
- Port scanning

Dataset:

- Total flows: 1,340
- Normal: 806 (60%)
- Attack: 534 (40%)



GET <http://localhost:8080/wm/core/switch/all/flow/json>

```
1 "flows": [
2     {
3         "match": {
4             "in_port": 1,
5             "eth_type": "0x0800",          // IPv4
6             "ipv4_src": "10.0.0.1",
7             "ipv4_dst": "10.0.0.2",
8             "ip_proto": "6",            // TCP
9             "tcp_src": 54321,
10            "tcp_dst": 80
11        },
12        "actions": [
13            {
14                "type": "OUTPUT",
15                "port": 2
16            }
17        ],
18        "packet_count": 42,           // <- Key statistic
19        "byte_count": 2688,          // <- Key statistic
20        "duration_sec": 5,          // <- Key statistic
21        "priority": 100,
22        "idle_timeout": 60,
23        "hard_timeout": 0,
24        "cookie": "0x0"
25    }
26    // ... more flows
27 ]
```

Features & Model

Raw Features:

- packet_count
- byte_count
- duration_sec

Derived Features:

- packet_rate (pkt/sec)
- byte_rate (bytes/sec)
- avg_packet_size

Training:

- 80/20 train/test split
- 30 epochs, Adam optimizer
- Cross-entropy loss

Model Architecture:

- LSTM (Long Short-Term Memory)
- Input: Sequences of 10 consecutive flows
- 2 layers, 64 hidden units each
- Output: Binary classification (normal/attack)

Results

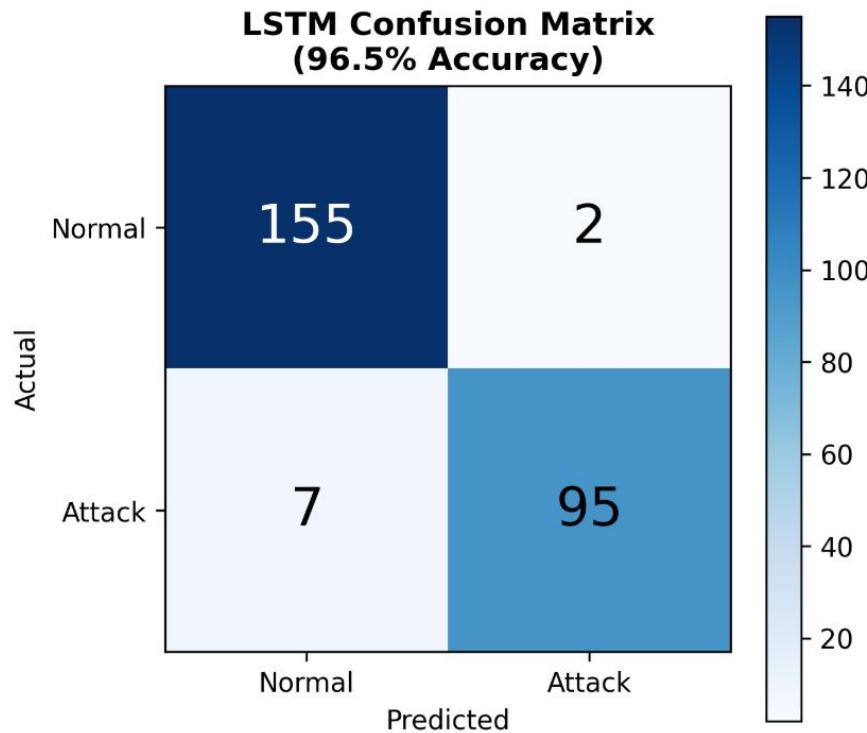
Model Performance:

Accuracy: 96.53%

Precision: 97.94%

Recall: 93.14%

F1-Score: 95.48%



Demo Script

```
=====
Normal Traffic
Flow: 10.0.0.2 → 10.0.0.5 [s1]
-----
Packets:      2 | Bytes:       108
Duration:   0.15s | Pkt Rate:  13.07 pkt/s
-----
→ NORMAL (100%) - ALLOW
=====

=====
Normal Traffic
Flow: 10.0.0.2 → 10.0.0.5 [s1]
-----
Packets:      2 | Bytes:       108
Duration:   0.15s | Pkt Rate:  13.61 pkt/s
-----
→ NORMAL (100%) - ALLOW
=====
```

```
=====
ATTACK DETECTED!
Flow: 10.0.0.5 → 10.0.0.2 [s1]
-----
Packets:      2 | Bytes:       108
Duration:   1.27s | Pkt Rate:  1.57 pkt/s
-----
→ ATTACK (99%) - BLOCK
=====

=====
ATTACK DETECTED!
Flow: 10.0.0.5 → 10.0.0.2 [s1]
-----
Packets:      2 | Bytes:       108
Duration:   1.27s | Pkt Rate:  1.57 pkt/s
-----
→ ATTACK (99%) - BLOCK
=====
```

```
=====
DEMO COMPLETE - FINAL RESULTS
=====
Total Flows:      27871
Normal:          2117
Attacks Detected: 25754
=====
```

Limitations

Current Limitations:

1. Simulated Environment
2. Limited Attack Diversity
3. Dataset Size
 - » 1,340 flows (proof-of-concept)
4. Flow Aggregation
 - » OpenFlow aggregates packets into flows

Future Work

- Test on real network traces
- Add more attack types (DDoS, SQL injection, etc.)
- Larger, more diverse dataset
- Deploy on production SDN controller
- Real-time anomaly response (block malicious flows)

Conclusion

- ✓ Proposed flow-based anomaly detection for SDN
- ✓ Used LSTM to learn from OpenFlow statistics
- ✓ Achieved 96.5% accuracy without external IDS
- ✓ Demonstrated feasibility of ML/DL-based SDN security

Key Contributions:

- Real-time flow-based detection
- Proof-of-concept validation